

HIVE ASSIGNMENT :USE OF SQL SCRIPT

SQL SCRIPT FILE

```
show databases;
create database acpbatchB;
use acpbatchB;
create table employee(name string, company string, sal int) row format
delimited fields terminated by '\t';
load data local inpath '/home/ccoew/ACP/employeebatchB.txt' into table
employee;
select * from employee;
select count(*) from employee;
select min(sal) from employee;
```

Input file

```
sarika      acb  10000
pradnya     hnm  2000
jaya frd    5000
radhatre    6000
```

Terminal output

```
ccoew@ccoew-System-Product-Name:~$ sudo su
[sudo] password for ccoew:
root@ccoew-System-Product-Name:/home/ccoew# start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
Starting namenodes on [localhost]
localhost: namenode running as process 13914. Stop it first.
localhost: datanode running as process 14057. Stop it first.
Starting secondary namenodes [0.0.0.0]
0.0.0.0: secondarynamenode running as process 14237. Stop it first.
starting yarn daemons
resourcemanager running as process 14392. Stop it first.
localhost: nodemanager running as process 14685. Stop it first.
root@ccoew-System-Product-Name:/home/ccoew# jps
13914 NameNode
17999 Jps
14237 SecondaryNameNode
```

14392 ResourceManager

14057 DataNode

14685 NodeManager

**root@ccoew-System-Product-Name:/home/ccoew# hive -f
/home/ccoew/ACP/batchbScript.sql**

16/09/29 16:50:16 INFO Configuration.deprecation:

mapred.input.dir.recursive is deprecated. Instead, use

mapreduce.input.fileinputformat.input.dir.recursive

16/09/29 16:50:16 INFO Configuration.deprecation: mapred.max.split.size

is deprecated. Instead, use mapreduce.input.fileinputformat.split.maxsize

16/09/29 16:50:16 INFO Configuration.deprecation: mapred.min.split.size is

deprecated. Instead, use mapreduce.input.fileinputformat.split.minsize

16/09/29 16:50:16 INFO Configuration.deprecation:

mapred.min.split.size.per.rack is deprecated. Instead, use

mapreduce.input.fileinputformat.split.minsize.per.rack

16/09/29 16:50:16 INFO Configuration.deprecation:

mapred.min.split.size.per.node is deprecated. Instead, use

mapreduce.input.fileinputformat.split.minsize.per.node

16/09/29 16:50:16 INFO Configuration.deprecation: mapred.reduce.tasks is

deprecated. Instead, use mapreduce.job.reduces

16/09/29 16:50:16 INFO Configuration.deprecation:

mapred.reduce.tasks.speculative.execution is deprecated. Instead, use

mapreduce.reduce.speculative

16/09/29 16:50:16 WARN conf.Configuration:

org.apache.hadoop.hive.conf.LoopingByteArrayInputStream@14332f47:an
attempt to override final parameter: mapreduce.job.end-
notification.max.retry.interval; Ignoring.

16/09/29 16:50:16 WARN conf.Configuration:

org.apache.hadoop.hive.conf.LoopingByteArrayInputStream@14332f47:an
attempt to override final parameter: mapreduce.job.end-
notification.max.attempts; Ignoring.

Logging initialized using configuration in jar:file:/usr/local/hive-

0.11.0/lib/hive-common-0.11.0.jar!/hive-log4j.properties

Hive history file=/tmp/root/hive_job_log_root_18444@ccoew-System-

Product-Name_201609291650_2124819542.txt

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in

[jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-

1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/usr/local/hive-0.11.0/lib/slf4j-log4j12-1.6.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]

OK

batchb

data4916

db001

db01

db1

db2

db3

db4

db40

db401

db402

db403

db61

db9

db91

db92

db93

db94

db95

db96

dbcl1

default

emp_records

empdb

empdb1

empdbc1

iamsuperhero4917

iamsuperhero4917_2

iamsuperhero4917_3

iamsuperhero4917_4

placement1

placement2

placement20

stud441

student1

zz1

Time taken: 2.747 seconds, Fetched: 36 row(s)

OK

Time taken: 0.449 seconds

OK

Time taken: 0.009 seconds

OK

Time taken: 0.503 seconds

Copying data from file:/home/ccoew/ACP/employeebatchB.txt

Copying file: file:/home/ccoew/ACP/employeebatchB.txt

Loading data to table acpbatchb.employee

Table acpbatchb.employee stats: [num_partitions: 0, num_files: 1, num_rows: 0, total_size: 63, raw_data_size: 0]

OK

Time taken: 0.54 seconds

OK

sarika acb 10000

pradnya hnm 2000

jaya frd 5000

radha tre 6000

Time taken: 0.19 seconds, Fetched: 4 row(s)

Total MapReduce jobs = 1

Launching Job 1 out of 1

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapred.reduce.tasks=<number>

16/09/29 16:50:22 INFO Configuration.deprecation:

mapred.input.dir.recursive is deprecated. Instead, use

mapreduce.input.fileinputformat.input.dir.recursive

16/09/29 16:50:22 INFO Configuration.deprecation: mapred.max.split.size

is deprecated. Instead, use mapreduce.input.fileinputformat.split.maxsize

16/09/29 16:50:22 INFO Configuration.deprecation: mapred.min.split.size is

deprecated. Instead, use mapreduce.input.fileinputformat.split.minsize

16/09/29 16:50:22 INFO Configuration.deprecation:
mapred.min.split.size.per.rack is deprecated. Instead, use
mapreduce.input.fileinputformat.split.minsize.per.rack
16/09/29 16:50:22 INFO Configuration.deprecation:
mapred.min.split.size.per.node is deprecated. Instead, use
mapreduce.input.fileinputformat.split.minsize.per.node
16/09/29 16:50:22 INFO Configuration.deprecation: mapred.reduce.tasks is
deprecated. Instead, use mapreduce.job.reduces
16/09/29 16:50:22 INFO Configuration.deprecation:
mapred.reduce.tasks.speculative.execution is deprecated. Instead, use
mapreduce.reduce.speculative
16/09/29 16:50:22 WARN conf.Configuration:
org.apache.hadoop.hive.conf.LoopingByteArrayInputStream@3410cdeb:an
attempt to override final parameter: mapreduce.job.end-
notification.max.retry.interval; Ignoring.
16/09/29 16:50:22 WARN conf.Configuration:
org.apache.hadoop.hive.conf.LoopingByteArrayInputStream@3410cdeb:an
attempt to override final parameter: mapreduce.job.end-
notification.max.attempts; Ignoring.
Execution log at: /tmp/root/.log
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in
[jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-
1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hive-0.11.0/lib/slf4j-log4j12-
1.6.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an
explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Job running in-process (local Hadoop)
Hadoop job information for null: number of mappers: 0; number of
reducers: 0
2016-09-29 16:50:24,993 null map = 100%, reduce = 100%
Ended Job = job_local376839168_0001
Execution completed successfully
Mapred Local Task Succeeded . Convert the Join into MapJoin
OK
4
Time taken: 4.448 seconds, Fetched: 1 row(s)
Total MapReduce jobs = 1

Launching Job 1 out of 1

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapred.reduce.tasks=<number>

16/09/29 16:50:26 INFO Configuration.deprecation:

mapred.input.dir.recursive is deprecated. Instead, use

mapreduce.input.fileinputformat.input.dir.recursive

16/09/29 16:50:26 INFO Configuration.deprecation: mapred.max.split.size

is deprecated. Instead, use mapreduce.input.fileinputformat.split.maxsize

16/09/29 16:50:26 INFO Configuration.deprecation: mapred.min.split.size is

deprecated. Instead, use mapreduce.input.fileinputformat.split.minsize

16/09/29 16:50:26 INFO Configuration.deprecation:

mapred.min.split.size.per.rack is deprecated. Instead, use

mapreduce.input.fileinputformat.split.minsize.per.rack

16/09/29 16:50:26 INFO Configuration.deprecation:

mapred.min.split.size.per.node is deprecated. Instead, use

mapreduce.input.fileinputformat.split.minsize.per.node

16/09/29 16:50:26 INFO Configuration.deprecation: mapred.reduce.tasks is

deprecated. Instead, use mapreduce.job.reduces

16/09/29 16:50:26 INFO Configuration.deprecation:

mapred.reduce.tasks.speculative.execution is deprecated. Instead, use

mapreduce.reduce.speculative

16/09/29 16:50:26 WARN conf.Configuration:

org.apache.hadoop.hive.conf.LoopingByteArrayInputStream@429a9a6e:an

attempt to override final parameter: mapreduce.job.end-

notification.max.retry.interval; Ignoring.

16/09/29 16:50:26 WARN conf.Configuration:

org.apache.hadoop.hive.conf.LoopingByteArrayInputStream@429a9a6e:an

attempt to override final parameter: mapreduce.job.end-

notification.max.attempts; Ignoring.

Execution log at: /tmp/root/.log

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in

[jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-

1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/usr/local/hive-0.11.0/lib/slf4j-log4j12-1.6.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]

Job running in-process (local Hadoop)

Hadoop job information for null: number of mappers: 0; number of reducers: 0

2016-09-29 16:50:29,070 null map = 100%, reduce = 100%

Ended Job = job_local937932227_0001

Execution completed successfully

Mapred Local Task Succeeded . Convert the Join into MapJoin

OK

2000

Time taken: 4.079 seconds, Fetched: 1 row(s)

root@ccoew-System-Product-Name:/home/ccoew#