

## MAP-REDUCE ASSIGNMENT

Checking frequency of user login through particular IP address with input log file

---

### Mapper class :Mpr

```
import java.io.IOException;
import java.util.regex.Matcher;
import java.util.regex.Pattern;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class Mpr extends Mapper<LongWritable, Text, Text, IntWritable> {

    private static final IntWritable one=new IntWritable(1);
    private Pattern patternIP=Pattern.compile("(?:(?:[0-9]+(?:\\.[0-9]+(?:\\.[0-9]+(?:[0-9]+)?)?)?)?)");
    private Pattern
patternURL=Pattern.compile("((https?|ftp|gopher|telnet|file):(((/)|(\\\\\\\\))+[\\w\\\\
d:#@%/$()~_?\\+\\-\\\\\\\\\\\\\\\\.&]*))");
    private Text word=new Text();
    public void map(LongWritable ikey, Text ivalue, Context context)
        throws IOException, InterruptedException {
        String url=null;
        String iPAddress=null;
        String[] entries=ivalue.toString().split("\\r?\\n");
        for (String string : entries) {
            Matcher matcherIP=patternIP.matcher(string);
            Matcher matcherUrl=patternURL.matcher(string);
            if(matcherIP.find())
            {
                iPAddress=matcherIP.group(1);
            }
            if(matcherUrl.find())
            {
                url=matcherUrl.group(1);
            }
        }
    }
}
```

```

        word.set(iPAddress+": "+url);
        context.write(word,one);
    }
}

```

### **Reducer class: Redr**

```

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class Redr extends Reducer<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text _key, Iterable<IntWritable> values, Context
context)
        throws IOException, InterruptedException {
        // process values
        int count=0;
        for (IntWritable val : values) {
            count += val.get();
        }
        context.write(_key, new IntWritable(count));
    }
}

```

## Driver class: Dri

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class Dri {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "College");
        job.setJarByClass(Dri.class);
        job.setMapperClass(Mpr.class);

        job.setReducerClass(Redr.class);

        // TODO: specify output types
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        // TODO: specify input and output DIRECTORIES (not files)
        FileInputFormat.setInputPaths(job, new Path(args[1]));
        FileOutputFormat.setOutputPath(job, new Path(args[2]));

        if (!job.waitForCompletion(true))
            return;
    }
}
```

**Input file : log.txt**

1	4	0	172.16.3.101	Sat Mar 14 13:09:20 IST 2015	Login
1	4	0	172.16.2.86	Sat Mar 14 15:07:14 IST 2015	Login
1	4	0	172.16.2.86	Sat Mar 14 15:10:48 IST 2015	Login
3	4	0	172.16.2.86	Sat Mar 14 15:12:07 IST 2015	Login
3	5	0	172.16.2.86	Sat Mar 14 15:13:08 IST 2015	
			USER-3/hadoop-2.6.0.tar.gz		
3	4	0	172.16.2.86	Sat Mar 14 15:13:16 IST 2015	Login
4	4	0	172.16.2.86	Sat Mar 14 15:16:46 IST 2015	Login
4	5	0	172.16.2.86	Sat Mar 14 15:17:02 IST 2015	
			USER-4/tokeHarami		
3	4	0	172.16.2.86	Sat Mar 14 15:17:25 IST 2015	Login
4	4	0	172.16.2.86	Sat Mar 14 15:19:02 IST 2015	Login
1	4	0	172.16.2.86	Sat Mar 14 15:20:42 IST 2015	Login
1	4	0	172.16.2.86	Sat Mar 14 15:25:24 IST 2015	Login
3	4	0	172.16.2.86	Sat Mar 14 15:25:44 IST 2015	Login
3	3	0	172.16.2.86	Sat Mar 14 15:26:24 IST 2015	ASD
3	4	0	172.16.2.86	Sat Mar 14 15:30:23 IST 2015	Login
3	4	0	172.16.2.86	Sat Mar 14 15:30:59 IST 2015	Login
3	4	0	172.16.2.86	Sat Mar 14 15:31:24 IST 2015	Login
5	4	0	172.16.2.110	Mon Mar 16 10:41:06 IST 2015	Login
4	4	0	172.16.2.110	Mon Mar 16 10:45:44 IST 2015	Login
4	5	0	172.16.2.110	Mon Mar 16 10:46:14 IST 2015	
			USER-4/StarToken-NG.msi		
4	4	0	172.16.2.110	Mon Mar 16 10:46:38 IST 2015	Login
5	4	0	172.16.2.110	Mon Mar 16 10:53:57 IST 2015	Login
5	3	0	172.16.2.110	Mon Mar 16 10:54:03 IST 2015	
			CHETAN		
5	3	0	172.16.2.110	Mon Mar 16 10:56:11 IST 2015	
			ABCD		
5	5	0	172.16.2.110	Mon Mar 16 11:00:44 IST 2015	
			USER-5/New Microsoft Office Excel Worksheet (2).xlsx		
5	5	0	172.16.2.110	Mon Mar 16 11:01:25 IST 2015	
			USER-5/New Microsoft Office Excel Worksheet (2).xlsx		
5	5	0	172.16.2.110	Mon Mar 16 11:01:35 IST 2015	
			USER-5/Class room pc status ( RAJU SM).docx		
5	5	0	172.16.2.110	Mon Mar 16 11:01:46 IST 2015	
			USER-5/Class room pc status ( RAJU SM).docx		

5	5	0	172.16.2.110	Mon Mar 16 11:02:14 IST 2015	
	USER-5/Hadoop-Introduction.pdf				
5	3	0	172.16.2.110	Mon Mar 16 11:09:26 IST 2015	
	ABCD				
6	4	0	172.16.2.43	Mon Mar 16 12:38:04 IST 2015	Login
6	5	0	172.16.2.43	Mon Mar 16 12:39:41 IST 2015	
	USER-6/cal7.lisp				
4	4	0	172.16.2.102	Tue Mar 17 11:13:51 IST 2015	Login
4	1	0	172.16.2.102	Tue Mar 17 11:14:39 IST 2015	
	USER-4/StarToken-NG.msi				
4	4	0	172.16.2.102	Tue Mar 17 11:15:58 IST 2015	Login
4	4	0	172.16.2.102	Tue Mar 17 11:17:04 IST 2015	Login
3	4	0	172.16.2.33	Fri Jun 19 11:03:04 IST 2015	Login
3	4	0	172.16.2.33	Fri Jun 19 11:03:40 IST 2015	Login
3	4	0	172.16.2.33	Fri Jun 19 11:04:33 IST 2015	Login
3	4	0	172.16.2.33	Fri Jun 19 11:04:54 IST 2015	Login
3	4	0	172.16.2.33	Fri Jun 19 11:08:17 IST 2015	Login
3	1	0	172.16.2.33	Fri Jun 19 11:08:37 IST 2015	
	USER-3/SMART TEST.docx				
3	3	0	172.16.2.33	Fri Jun 19 11:09:10 IST 2015	C
3	1	0	172.16.2.33	Fri Jun 19 11:09:16 IST 2015	
	/SharedData/4/C programming 1st_edition.pdf				
3	3	0	172.16.2.33	Fri Jun 19 11:13:10 IST 2015	
	CSYLLABUS				
3	3	0	172.16.2.33	Fri Jun 19 11:13:14 IST 2015	
	YLLABUS				
3	1	0	172.16.2.33	Fri Jun 19 11:13:17 IST 2015	
	/SharedData/2/Syllabus TE Computer Engineering 2012Course.pdf				
3	3	0	172.16.2.33	Fri Jun 19 11:14:44 IST 2015	C
3	1	0	172.16.2.33	Fri Jun 19 11:14:58 IST 2015	
	/SharedData/4/C programming 1st_edition.pdf				
3	3	0	172.16.2.33	Fri Jun 19 11:15:43 IST 2015	
	YLLABUS				
3	3	0	172.16.2.33	Fri Jun 19 11:15:49 IST 2015	
	SYLLABUS				
3	3	0	172.16.2.33	Fri Jun 19 11:15:53 IST 2015	BUS
5	3	0	172.16.2.110	Wed Jul 08 15:15:51 IST 2015	
	DEMO				
5	3	0	172.16.2.110	Wed Jul 08 15:15:51 IST 2015	
	DEMO				

5	3	0	172.16.2.110	Wed Jul 08 15:15:51 IST 2015	
			DEMO		
5	3	0	172.16.2.110	Wed Jul 08 15:15:51 IST 2015	
			DEMO		
5	3	0	172.16.2.110	Wed Jul 08 15:15:51 IST 2015	
			DEMO		
5	3	0	172.16.2.110	Wed Jul 08 15:15:52 IST 2015	
			DEMO		
5	4	0	172.16.2.110	Wed Jul 08 15:21:45 IST 2015	Login
5	4	0	172.16.2.110	Wed Jul 08 15:24:44 IST 2015	Login
5	3	0	172.16.2.110	Wed Jul 08 15:27:15 IST 2015	A
3	4	0	172.16.3.100	Wed Jul 08 15:32:05 IST 2015	Login

### Output file: part-r-00000

172.16.1.124:null	2
172.16.1.125:null	15
172.16.1.144:null	7
172.16.1.177:null	17
172.16.1.178:null	9
172.16.1.182:null	3
172.16.1.183:null	4
172.16.1.184:null	5
172.16.1.185:null	9
172.16.1.186:null	7
172.16.1.187:null	12
172.16.1.188:null	45
172.16.1.189:null	16
172.16.1.190:null	13
172.16.1.193:null	55
172.16.1.195:null	23
172.16.1.198:null	18
172.16.1.199:null	1
172.16.1.209:null	2
172.16.1.217:null	9
172.16.1.222:null	9
172.16.2.58:null	4

172.16.2.59:null	7
172.16.2.60:null	8
172.16.2.62:null	4
172.16.2.64:null	1
172.16.2.69:null	1
172.16.2.86:null	79
172.16.2.87:null	8
172.16.2.88:null	51
172.16.2.89:null	5
172.16.2.8:null	1
172.16.2.90:null	21
172.16.2.92:null	46
172.16.2.94:null	17
172.16.2.95:null	5
172.16.2.96:null	11
172.16.2.97:null	23
172.16.2.99:null	14
172.16.3.100:null	3
172.16.3.101:null	187
172.16.3.104:null	12
172.16.3.105:null	6
172.16.3.108:null	1

Terminal commands to start hadoop

***root@ccoew-desktop:/usr/local/hadoop/etc/hadoop# start-dfs.sh***

Starting namenodes on [localhost]

localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-root-namenode-ccoew-desktop.out

localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-root-datanode-ccoew-desktop.out

Starting secondary namenodes [0.0.0.0]

The authenticity of host '0.0.0.0 (0.0.0.0)' can't be established.

ECDSA key fingerprint is fc:7a:18:cf:cc:46:85:73:c0:1b:5e:fc:ac:bc:c8:01.

**Are you sure you want to continue connecting (yes/no)? yes**

0.0.0.0: Warning: Permanently added '0.0.0.0' (ECDSA) to the list of known hosts.

0.0.0.0: starting secondarynamenode, logging to

/usr/local/hadoop/logs/hadoop-root-secondarynamenode-ccoew-desktop.out

***root@ccoew-desktop:/usr/local/hadoop/etc/hadoop# start-yarn.sh***

starting yarn daemons

starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-root-resourcemanager-ccoew-desktop.out

***root@ccoew-desktop:/usr/local/hadoop/etc/hadoop# jps***

3559 SecondaryNameNode

4099 Jps

3388 DataNode

3243 NameNode

3699 ResourceManager

3816 NodeManager

***root@ccoew-desktop:/usr/local/hadoop/etc/hadoop#***

Terminal commands for mapreduce :

***root@ccoew-desktop:/home/ccoew# hdfs dfs -mkdir /ankita***

***root@ccoew-desktop:/home/ccoew# hdfs dfs -put  
/home/ccoew/log.txt /ankita***

***root@ccoew-desktop:/home/ccoew# hadoop jar /home/ccoew/p1.jar  
Dri /ankita/log.txt /ankita/op***

16/09/23 16:53:54 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id

16/09/23 16:53:54 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=

16/09/23 16:53:54 WARN mapreduce.JobSubmitter: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.

16/09/23 16:53:54 INFO input.FileInputFormat: Total input paths to process : 1

16/09/23 16:53:54 INFO mapreduce.JobSubmitter: number of splits:1

16/09/23 16:53:54 INFO mapreduce.JobSubmitter: Submitting tokens for job: job\_local949013430\_0001

16/09/23 16:53:55 INFO mapreduce.Job: The url to track the job: http://localhost:8080/



16/09/23 16:53:55 INFO mapreduce.Job: Running job:  
job\_local949013430\_0001  
16/09/23 16:53:55 INFO mapred.LocalJobRunner: OutputCommitter set in  
config null  
16/09/23 16:53:55 INFO mapred.LocalJobRunner: OutputCommitter is  
org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter  
16/09/23 16:53:55 INFO mapred.LocalJobRunner: Waiting for map tasks  
16/09/23 16:53:55 INFO mapred.LocalJobRunner: Starting task:  
attempt\_local949013430\_0001\_m\_000000\_0  
16/09/23 16:53:55 INFO mapred.Task: Using  
ResourceCalculatorProcessTree : [ ]  
16/09/23 16:53:55 INFO mapred.MapTask: Processing split:  
hdfs://localhost:9000/ankita/log.txt:0+124829  
16/09/23 16:53:55 INFO mapred.MapTask: (EQUATOR) 0 kvi  
26214396(104857584)  
16/09/23 16:53:55 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100  
16/09/23 16:53:55 INFO mapred.MapTask: soft limit at 83886080  
16/09/23 16:53:55 INFO mapred.MapTask: bufstart = 0; bufvoid =  
104857600  
16/09/23 16:53:55 INFO mapred.MapTask: kvstart = 26214396; length =  
6553600  
16/09/23 16:53:55 INFO mapred.MapTask: Map output collector class =  
org.apache.hadoop.mapred.MapTask\$MapOutputBuffer  
16/09/23 16:53:55 INFO mapred.LocalJobRunner:  
16/09/23 16:53:55 INFO mapred.MapTask: Starting flush of map output  
16/09/23 16:53:55 INFO mapred.MapTask: Spilling map output  
16/09/23 16:53:55 INFO mapred.MapTask: bufstart = 0; bufend = 40241;  
bufvoid = 104857600  
16/09/23 16:53:55 INFO mapred.MapTask: kvstart =  
26214396(104857584); kvend = 26206908(104827632); length =  
7489/6553600  
16/09/23 16:53:55 INFO mapred.MapTask: Finished spill 0  
16/09/23 16:53:55 INFO mapred.Task:  
Task:attempt\_local949013430\_0001\_m\_000000\_0 is done. And is in the  
process of committing  
16/09/23 16:53:55 INFO mapred.LocalJobRunner: map  
16/09/23 16:53:55 INFO mapred.Task: Task  
'attempt\_local949013430\_0001\_m\_000000\_0' done.  
16/09/23 16:53:55 INFO mapred.LocalJobRunner: Finishing task:  
attempt\_local949013430\_0001\_m\_000000\_0

16/09/23 16:53:55 INFO mapred.LocalJobRunner: map task executor complete.

16/09/23 16:53:55 INFO mapred.LocalJobRunner: Waiting for reduce tasks

16/09/23 16:53:55 INFO mapred.LocalJobRunner: Starting task:  
attempt\_local949013430\_0001\_r\_000000\_0

16/09/23 16:53:55 INFO mapred.Task: Using  
ResourceCalculatorProcessTree : [ ]

16/09/23 16:53:55 INFO mapred.ReduceTask: Using  
ShuffleConsumerPlugin:  
org.apache.hadoop.mapreduce.task.reduce.Shuffle@7bfbde18

16/09/23 16:53:55 INFO reduce.MergeManagerImpl: MergerManager:  
memoryLimit=333971456, maxSingleShuffleLimit=83492864,  
mergeThreshold=220421168, ioSortFactor=10,  
memToMemMergeOutputsThreshold=10

16/09/23 16:53:55 INFO reduce.EventFetcher:  
attempt\_local949013430\_0001\_r\_000000\_0 Thread started: EventFetcher  
for fetching Map Completion Events

16/09/23 16:53:55 INFO reduce.LocalFetcher: localfetcher#1 about to  
shuffle output of map attempt\_local949013430\_0001\_m\_000000\_0  
decomp: 43989 len: 43993 to MEMORY

16/09/23 16:53:55 INFO reduce.InMemoryMapOutput: Read 43989 bytes  
from map-output for attempt\_local949013430\_0001\_m\_000000\_0

16/09/23 16:53:55 INFO reduce.MergeManagerImpl: closeInMemoryFile ->  
map-output of size: 43989, inMemoryMapOutputs.size() -> 1,  
commitMemory -> 0, usedMemory -> 43989

16/09/23 16:53:55 INFO reduce.EventFetcher: EventFetcher is interrupted..  
Returning

16/09/23 16:53:55 INFO mapred.LocalJobRunner: 1 / 1 copied.

16/09/23 16:53:55 INFO reduce.MergeManagerImpl: finalMerge called with  
1 in-memory map-outputs and 0 on-disk map-outputs

16/09/23 16:53:55 INFO mapred.Merger: Merging 1 sorted segments

16/09/23 16:53:55 INFO mapred.Merger: Down to the last merge-pass,  
with 1 segments left of total size: 43969 bytes

16/09/23 16:53:55 INFO reduce.MergeManagerImpl: Merged 1 segments,  
43989 bytes to disk to satisfy reduce memory limit

16/09/23 16:53:55 INFO reduce.MergeManagerImpl: Merging 1 files, 43993  
bytes from disk

16/09/23 16:53:55 INFO reduce.MergeManagerImpl: Merging 0 segments,  
0 bytes from memory into reduce

16/09/23 16:53:55 INFO mapred.Merger: Merging 1 sorted segments

16/09/23 16:53:55 INFO mapred.Merger: Down to the last merge-pass,  
with 1 segments left of total size: 43969 bytes  
16/09/23 16:53:55 INFO mapred.LocalJobRunner: 1 / 1 copied.  
16/09/23 16:53:55 INFO Configuration.deprecation: mapred.skip.on is  
deprecated. Instead, use mapreduce.job.skiprecords  
16/09/23 16:53:56 INFO mapreduce.Job: Job job\_local949013430\_0001  
running in uber mode : false  
16/09/23 16:53:56 INFO mapreduce.Job: map 100% reduce 0%  
16/09/23 16:53:56 INFO mapred.Task:  
Task:attempt\_local949013430\_0001\_r\_000000\_0 is done. And is in the  
process of committing  
16/09/23 16:53:56 INFO mapred.LocalJobRunner: 1 / 1 copied.  
16/09/23 16:53:56 INFO mapred.Task: Task  
attempt\_local949013430\_0001\_r\_000000\_0 is allowed to commit now  
16/09/23 16:53:56 INFO output.FileOutputCommitter: Saved output of task  
'attempt\_local949013430\_0001\_r\_000000\_0' to  
hdfs://localhost:9000/ankita/op/\_temporary/0/task\_local949013430\_0001\_r\_000000\_0  
16/09/23 16:53:56 INFO mapred.LocalJobRunner: reduce > reduce  
16/09/23 16:53:56 INFO mapred.Task: Task  
'attempt\_local949013430\_0001\_r\_000000\_0' done.  
16/09/23 16:53:56 INFO mapred.LocalJobRunner: Finishing task:  
attempt\_local949013430\_0001\_r\_000000\_0  
16/09/23 16:53:56 INFO mapred.LocalJobRunner: reduce task executor  
complete.  
16/09/23 16:53:57 INFO mapreduce.Job: map 100% reduce 100%  
16/09/23 16:53:57 INFO mapreduce.Job: Job job\_local949013430\_0001  
completed successfully  
16/09/23 16:53:57 INFO mapreduce.Job: Counters: 38

#### File System Counters

FILE: Number of bytes read=99538  
FILE: Number of bytes written=642043  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=249658  
HDFS: Number of bytes written=1689  
HDFS: Number of read operations=13  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=4

## Map-Reduce Framework

Map input records=1873  
Map output records=1873  
Map output bytes=40241  
Map output materialized bytes=43993  
Input split bytes=101  
Combine input records=0  
Combine output records=0  
Reduce input groups=84  
Reduce shuffle bytes=43993  
Reduce input records=1873  
Reduce output records=84  
Spilled Records=3746  
Shuffled Maps =1  
Failed Shuffles=0  
Merged Map outputs=1  
GC time elapsed (ms)=28  
CPU time spent (ms)=0  
Physical memory (bytes) snapshot=0  
Virtual memory (bytes) snapshot=0  
Total committed heap usage (bytes)=496500736

## Shuffle Errors

BAD\_ID=0  
CONNECTION=0  
IO\_ERROR=0  
WRONG\_LENGTH=0  
WRONG\_MAP=0  
WRONG\_REDUCE=0

## File Input Format Counters

Bytes Read=124829

## File Output Format Counters

Bytes Written=1689

***root@ccoew-desktop:/home/ccoew#***